

Final Project

● Graded

Student

Haoyu Zhang

Total Points

20 / 20 pts

Autograder Score

0.0 / 0.0

Question 2

Final Project

20 / 20 pts

✓ - 0 pts Excellent

Autograder Results

This assignment does not have an autograder configured.

Submitted Files

SMCI Stock Predictor Project

Author: Haoyu Zhang

Course Project, UC Irvine, Math 10, Spring 25

I would like to post my notebook on the course's website. [Yes](#)

I) Introduction

The purpose of this project is to predict the stock price of SMCI (Super Micro Computer, Inc.). SMCI has shown significant growth and volatility in recent years, which makes it an interesting candidate for time-series modeling. We will explore its historical price data, visualize patterns, engineer features, and use regression models (linear and tree-based) to predict the next-day closing price.

II) Importing Data

We import the required libraries and load SMCI's historical stock data from a manually provided CSV file (MAX range).

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score

df = pd.read_csv("HistoricalData_1749502320765.csv")
df['Close/Last'] = df['Close/Last'].str.replace('$', '', regex=False).astype(float)
df['Open'] = df['Open'].str.replace('$', '', regex=False).astype(float)
df['High'] = df['High'].str.replace('$', '', regex=False).astype(float)
df['Low'] = df['Low'].str.replace('$', '', regex=False).astype(float)
df['Volume'] = df['Volume'].astype(str).str.replace(',', '').astype(int)
df['Date'] = pd.to_datetime(df['Date'])
df.sort_values('Date', inplace=True)
df.head()
```

Out [1]:

	Date	Close/Last	Volume	Open	High	Low
2514	2015-06-09	3.344	2886240	3.363	3.379	3.321
2513	2015-06-10	3.465	5746580	3.353	3.481	3.353
2512	2015-06-11	3.474	6701280	3.447	3.529	3.436

```
2511 2015-06-12    3.393 4931020 3.450 3.467 3.385  
2510 2015-06-15    3.292 7650700 3.426 3.447 3.286
```

III) Sorting & Cleaning Data

Here we remove missing values and add rolling averages (technical indicators).

```
In [2]:  
df.dropna(inplace=True)  
df['MA_10'] = df['Close/Last'].rolling(window=10).mean()  
df['MA_50'] = df['Close/Last'].rolling(window=50).mean()  
df.head()
```

```
Out [2]:  
          Date Close/Last  Volume  Open  High  Low  MA_10  MA_50  
2514 2015-06-09    3.344 2886240 3.363 3.379 3.321   NaN   NaN  
2513 2015-06-10    3.465 5746580 3.353 3.481 3.353   NaN   NaN  
2512 2015-06-11    3.474 6701280 3.447 3.529 3.436   NaN   NaN  
2511 2015-06-12    3.393 4931020 3.450 3.467 3.385   NaN   NaN  
2510 2015-06-15    3.292 7650700 3.426 3.447 3.286   NaN   NaN
```

IV) Data Exploration

We examine general statistics, visualize stock trends, and analyze feature correlations.

```
In [3]: df.describe()
```

```
Out [3]:  
          Date Close/Last  Volume  Open  \  
count      2515 2515.000000 2.515000e+03 2515.000000  
mean  2020-06-04 21:46:35.546719488 12.246801 1.684573e+07 12.247913  
min   2015-06-09 00:00:00  1.165000 3.038000e+04  1.155000  
25%   2017-12-04 12:00:00  2.304500 2.327225e+06  2.301500  
50%   2020-06-05 00:00:00  2.894000 4.028380e+06  2.885000  
75%   2022-12-01 12:00:00  8.051500 1.252550e+07  8.040000  
max   2025-06-06 00:00:00 118.807000 3.697348e+08 121.200000  
std     NaN 20.981443 3.332872e+07 21.037977  
  
          High  Low  MA_10  MA_50  
count 2515.000000 2515.000000 2506.000000 2466.000000  
mean  12.665328 11.837546 12.209585 12.077759  
min   1.216000 0.850000 1.223300 1.383000  
25%   2.346500 2.269500 2.317613 2.313607  
50%   2.950000 2.835000 2.869550 2.777750  
75%   8.224800 7.859000 7.970850 7.719190  
max   122.900000 112.234000 112.198700 94.969060  
std   21.822282 20.171035 20.883752 20.516064
```

In [4]:

```
plt.figure(figsize=(12,6))
plt.plot(df['Date'], df['Close/Last'], label='Close')
plt.plot(df['Date'], df['MA_10'], label='10-day MA')
plt.plot(df['Date'], df['MA_50'], label='50-day MA')
plt.title('SMCI Stock Closing Price & Moving Averages')
plt.xlabel('Date')
plt.ylabel('Price')
plt.legend()
plt.grid(True)
plt.show()
```




In [5]:

```
plt.figure(figsize=(10,5))
sns.histplot(df['Close/Last'], bins=100, kde=True, color='steelblue')
plt.title('Distribution of SMCI Closing Prices')
plt.xlabel('Price (USD)')
plt.ylabel('Frequency')
plt.grid(True)
plt.show() # Histogram
```

/opt/anaconda3/lib/python3.11/site-packages/seaborn/_oldcore.py:1119: FutureWarning


```
with pd.option_context('mode.use_inf_as_na', True):
```



Histogram Analysis:

The SMCI closing prices are right-skewed, showing that the stock traded at relatively lower prices most of the time but experienced significant upward spikes. This suggests a volatile growth pattern and a few high-price outliers, especially in recent years.

```
In [6]:  
plt.figure(figsize=(8,6)) # Correlation heatmap  
sns.heatmap(df[['Open', 'High', 'Low', 'Close/Last', 'Volume']].corr(), annot=True,  
cmap='coolwarm')  
plt.title('Feature Correlation Heatmap')  
plt.show()
```



Heatmap Analysis:

There is strong positive correlation between Open, High, Low, and Close prices, which is expected in stock data. Volume, however, has weaker correlation, suggesting it may have less predictive power for price modeling.

V) Lagged Prices and Moving Averages

We engineer lagged features and moving averages to help the model learn from previous trends.

```
In [7]:  
for lag in [1, 2, 3]: # Lagged features  
    df[f'Lag_{lag}'] = df['Close/Last'].shift(lag)  
  
for ma in [5, 10, 20]: # Lagged features  
    df[f'MA_{ma}'] = df['Close/Last'].rolling(window=ma).mean()  
  
df.dropna(inplace=True)  
df.head()
```

Out [7]:

	Date	Close/Last	Volume	Open	High	Low	MA_10	MA_50
2465	2015-08-18	2.750	3332050	2.750	2.771	2.7450	2.7170	2.84324
2464	2015-08-19	2.744	3694050	2.738	2.765	2.6930	2.7123	2.83124
2463	2015-08-20	2.674	5048290	2.736	2.756	2.6660	2.7110	2.81542
2462	2015-08-21	2.539	6602320	2.637	2.683	2.5302	2.6991	2.79672
2461	2015-08-24	2.456	6735940	2.450	2.605	2.3440	2.6729	2.77798

	Lag_1	Lag_2	Lag_3	MA_5	MA_20
2465	2.765	2.755	2.728	2.7404	2.64695
2464	2.750	2.765	2.755	2.7484	2.65720
2463	2.744	2.750	2.765	2.7376	2.66210
2462	2.674	2.744	2.750	2.6944	2.66180
2461	2.539	2.674	2.744	2.6326	2.65860

VI) Linear Regression

We fit a linear model using key numerical features to predict the next-day price.

```
In [8]: df['Next_Close'] = df['Close/Last'].shift(-1) # Features and target
df.dropna(inplace=True)
features = ['Open', 'High', 'Low', 'Close/Last', 'Volume', 'Lag_1', 'MA_5', 'MA_10']
X = df[features]
y = df['Next_Close']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model_lr = LinearRegression()
model_lr.fit(X_train, y_train)
y_pred_lr = model_lr.predict(X_test)

print("Linear Regression R^2:", r2_score(y_test, y_pred_lr))
print("MSE:", mean_squared_error(y_test, y_pred_lr))
```

Linear Regression R²: 0.9946140397348965
MSE: 2.6390893816913406

VII) Tree-Based Model (Random Forest)

To capture nonlinear interactions, we apply a Random Forest Regressor.

```
In [9]: model_rf = RandomForestRegressor(n_estimators=100, random_state=42)
model_rf.fit(X_train, y_train)
y_pred_rf = model_rf.predict(X_test)
```

```
print("Random Forest R2:", r2_score(y_test, y_pred_rf))
print("MSE:", mean_squared_error(y_test, y_pred_rf))
```

Random Forest R²: 0.9939059099020344
MSE: 2.986068904520977

- $R^2 = 0.9939$

This means the model explains over 99% of the variance in the next-day SMCI closing price. An R^2 this high suggests an extremely strong fit, indicating that the model has captured the relationship between input features (price, volume, lags, MAs) and the target variable very well.

- $MSE \approx 2.99$

This is a low value in the context of stock prices (especially if average SMCI prices are much higher). It shows that on average, the predicted price is off by less than \$3, which is relatively minor for a volatile tech stock.

Random Forest performed significantly better than the linear regression model, highlighting its strength in handling nonlinear patterns and interactions between features. By incorporating lagged prices and moving averages, the model was able to learn from historical momentum and trends, boosting its predictive power. This result demonstrates the importance of: Feature engineering in time-series problems Using ensemble models for complex, noisy financial data

VIII) Discussion of Results

- The linear regression model provides a simple baseline, but underperforms on volatile days.
- The Random Forest model performs better, likely due to capturing nonlinear price dynamics and lagged features.
- Lag features and moving averages improve prediction quality.
- While a high R^2 is good, it might also signal potential overfitting. This means the model fits the training data too closely and may not generalize as well to unseen market conditions.
- Financial markets are influenced by external macroeconomic events (news, earnings reports, etc.) not captured in this dataset.

IX) Cross-Validation and Bias–Variance Tradeoff

In future work, K-Fold Cross Validation can help evaluate model stability.

- Linear models may underfit due to high bias.
- Tree-based models may overfit unless tuned, but tend to have lower bias.

X) Summary & Future Improvements

This project successfully applied EDA and regression to SMCI stock data. Future improvements may include:

- Cross-validation for robustness
- Hyperparameter tuning
- More advanced models (XGBoost, LSTM)
- Using macroeconomic indicators

XI) References

- Yahoo Finance: <https://finance.yahoo.com/quote/SMCI/history/>
- Referring to this sample project: https://rayzhangzirui.github.io/math10fa24/final_project
- Downloading the historical quotes datas: <https://www.nasdaq.com/market-activity/stock>
- Scikit-learn documentation: <https://scikit-learn.org/>
- Pandas Library Documentation: <https://pandas.pydata.org/>
- numpy Library Documentation: <https://numpy.org/>
- matplotlib Library Documentation: <https://matplotlib.org/>
- seaborn Library Documentation: <https://seaborn.pydata.org/>
- Tree-Based Models:
https://www.researchgate.net/publication/378435618_Stock_price_prediction_using_decision_trees
- Lecture notes by Professor (RayZirui) Zhang: <https://rayzhangzirui.github.io/math10fa24>